Updates & Common Errors

Outlier

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-03-06T18:07:57.899Z

Purpose ▪ Scope ▪ Structure

Purpose

Overview and clarification of common errors in order to improve quality alignment.

Scope

Project-wide common errors for Eunomia only. Any individual account issues should be addressed by reaching out to a QM.

Structure

Common Errors presentation with time for clarification questions.
Please keep your mic muted when not speaking. This helps everyone focus.
Clarification questions may be asked by raising your hand or adding the question to the Zoom chat.

CONFIDENTIAL | ruban.arnold@gmail.com+outlier | 2026-03-06T18:07:57.889Z

Taxonomy Update

The overall scoring and justification step covering all models has been removed.

Updated Workflow

- Evaluate the prompt, create rubrics, and grade the rubric criteria for each of the responses as usual.
- You will see the scoring sheet showing which criteria the responses passed or failed.
- Check the differences in scoring between the responses, but ***you will no longer have to assign overall scores or write justifications.***
- Select the highest scoring model response and use it as the base to write the golden response.
- This change will allow you to **dedicate more time to the Golden response rewrite to ensure that it passes all rubric criteria**.

Key Opportunity & Warning ⚠️

Submission of failing rewrites remains a frequent error. Utilize this new structure to dedicate more time and space for improving Golden Response quality!

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-03-06T18:07:57.889Z

Update

📢 Project Update: Enhanced Prompts & Rubric Strategy

We have uploaded a new batch of Multi-Turn Tasks (3–4 turns). To ensure we are differentiating models effectively, we are shifting our rubric strategy.

Please follow these guidelines to ensure your tasks meet the new quality bar while remaining safe for QC audits.

The Strategy: Focus on "Mandatory Fails"

Our goal is to move away from using subjective Quality of Life (QoL) polish to rank models. Instead, we **want to see models fail because they couldn't handle complex instructions**.

- Instruction Following (IF) = Always Mandatory.
- No "40% QoL" Rule: For these complex tasks, focus on high-impact IF and Code rubrics. You do not need to pad the task with QoL rubrics just to hit a percentage.
- Mandatory Weighting: Use Mandatory for anything essential to the user's intent. Use Nice-to-Have (NtH) for extra polish or optimizations.

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-03-06T18:07:57.889Z

Update

Explicit vs. Implicit Requirements

In multi-turn conversations, requirements can come from two places. To force model fails, you should look for both:

Type	Definition	Example Scenario	Correct Atomic & Self-Contained Rubric Example	QC Impact
Explicit	Stated in the **current prompt** or **clearly referenced** from the past.	**Prompt:** "Use the *second strategy* (Python) we discussed."	**IF [Mandatory]:** "The response implements the logic using **Python** …" [with continuation specifying the strategy without referencing the prompt history]	**QC FAIL** if missing.
Contextual Implicit	Constraints from **past turns** NOT specifically referenced in the current prompt.	**Turn 1:** "Use React." **Prompt** "Add a loop." (No mention of React in T4).	**IF [Mandatory]:** "The code is written using **React.js** components …" "The code includes a loop …"	**QC** Score capped at 3 if missing.
Logical Implicit	**Not stated**, but functionally necessary to complete the user's intent.	**Prompt:** "Build a personalized fitness tracker."	**IF [Mandatory]:** "The response includes a **user profile data structure** (e.g., age, weight, or goals) to enable personalization."	**QC** Score capped at 3 if missing.

All **Instruction Following (IF)** rubrics **must be weighted as Mandatory**.

CONFIDENTIAL | ruban.arnold@gmail.com+outlier | 2026-03-06T18:07:57.889Z

Update

Summary:

- Scan your history! If the user says "do it again," look back at what "it" was! That is now an **Explicit Mandatory** requirement. For everything else, use your technical expertise to add "Implicit" rubrics that force the models to prove their intelligence.
- **Implicit** requirements can come from previous conversation history, not just the prompt itself!

💡 **The Strategy:** We *encourage* you to write **Mandatory IF** rubrics for these logical Implicit requirements. This is the best way to catch a model "hallucinating" a half-baked solution.

Note that you do not have to flag explicit and implicit requirements separately, but knowing the difference can help clarify where to look for them and how to describe them in justification.

Remember that rubrics must never mention the prompt or previous turns. They must be self-contained so that the response could be evaluated based on rubrics alone without reference to the prompt, chat history, or outside resources.

CONFIDENTIAL | ruban.arnold@gmail.com+outlier | 2026-03-06T18:07:57.889Z

Update

Rubrics Missing Coverage

The rubrics should exhaustively and collectively check if all the explicit requirements and implicit expectations of the prompt are addressed.

- **QC FAIL (Missing Coverage):** Your task will Fail if you miss any **Explicit** requirement written in the prompt.
- **Penalty Reminder:** While missing an Implicit instruction is not a QC "Fail," it is **penalized**. Missing a critical **implicit** IF rubric will result in a maximum task score of 3.

Ratability in Multi-Turn Tasks

With 3–4 turns, do not judge a prompt in isolation.

- **Context is King:** If the final prompt is short (e.g., *"Now fix the bug"*), it is **Ratable** because it relies on the complex code generated in previous turns.
- **Complexity:** If the total conversation history provides a technical challenge, the task is valuable. Do not mark it as "Too Easy."
- **Niche:** Still need to watch out for niche platforms or products. They should be non-ratable.

CONFIDENTIAL | ruban.arnold@gmail.com+outlier | 2026-03-06T18:07:57.889Z

Update

Special Case: Code with No Instructions

When a user provides only code (no text command) and there is no past context:

- Identify the Intent: Is the code broken? Is it a snippet needing completion?
- Implicit IF: The "Instruction" is to provide a functional, debugged, or completed version of that code.
- The Rubric: Create Mandatory IF rubrics based on the logical goal

Example: Among multiple issues, the user provided a query with a syntax error near a join.

❌ **Bad Criterion** "The response successfully identifies and fixes the syntax error in the provided snippet" is not a self-contained rubric criterion because it references the chat context.

✅ **Good Criterion** “The SQL query follows standard execution syntax, placing the JOIN condition sequentially after the FROM clause and before the WHERE clause.”

Update

🚀 Summary Checklist

1. **Check the Chat History:** Scan all 3–4 turns for explicit constraints (language, libraries, logic).
2. **Add "Technical Teeth":** Look for explicit and implicit requirements to turn into **Mandatory IF** or Code rubrics.
3. **Add Nice-to-Have:** Add **NtH QoL** rubrics to further differentiate the models.
4. **Category Check:**
   - **IF:** Always Mandatory.
   - **Code:** Mandatory (Logic) or NtH (Optimization).
   - **QoL:** Always Nice-to-Have (NtH).

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-03-06T18:07:57.889Z

Update

Rubric Categories and Weights

Category	Weight	Logic
IF	**Mandatory**	Explicit asks + Critical Implicit intent.
Code	**Mandatory / NtH**	Technical logic and correctness.
QoL	**Nice-to-Have**	Visual style, tone, and formatting polish.

CONFIDENTIAL | ruban.arnold@gmail.com+outlier | 2026-03-06T18:07:57.899Z

Update

Avoid Misunderstandings of Number of Required Elements

We do not require 10 rubrics

- We recommend ~10 rubrics as general guideline
- Tasks with fewer than 10 rubric criteria should not be penalized based on the number of rubrics.
- As long as there are **enough criteria to cover the prompt requirements** in accordance with project specifications, it is not a problem to have more or fewer than 10 rubrics.
- We no longer have a 40% QoL guideline. QoL rubrics should be noted and present, but we are now more focused on mandatory Instruction Following rubrics that can force a failing score.

We have seen some reviewers leave feedback stating that there must be at least 10 rubric criteria per task or at least 40% QoL. Such feedback is not correct at this time.

CONFIDENTIAL | ruban.arnold@gmail.com+outlier | 2026-03-06T18:07:57.899Z

Rubric Category definitions

Code

→ Evaluates the **correctness, safety,** and **intent** of the code

- Focus on intent, not exact syntax
- Avoid requiring specific methods, libraries, or line-by-line implementations unless explicitly needed
- Accept alternative valid approaches if they meet the goal

Instruction Following (IF)

→ Ensures the model correctly follows the user's **explicit and implicit instructions**

- Check that all parts of the prompt are fully addressed
- No additions, omissions, or scope drift
- No reinterpreting or ignoring user requirements

Instruction Following (IF) rubrics must be mandatory.

Quality of Life (QoL)

→ Evaluates usefulness and clarity across three sub-dimensions:

- **Helpfulness** – relevance, usefulness, guidance, and solving the user's need
- **Style & Clarity** – readable code, clear structure, logical flow, visual organization
- **Verbosity** – concise yet complete; no rambling, no missing essential information

Quality of Life Rubrics (QoL) must be Nice-to-have.

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-03-06T18:07:57.889Z

Golden Response

Most Common QC Audit Error: Golden Response Rewrite Rewrite Failures

The most frequent error is that the Golden Response rewrite fails one or more rubric criteria.

When a Golden Response (rewrite) fails any of the rubric criteria, the maximum score it can receive is **2**.

How to Avoid This Error:

- ✔ Read in-task evals and resolve any failed rubric criteria before submitting.
- ✔ Remove filler phrases, unnecessary praise, and tone-deaf affirmations from the model response.
- ✔ Leverage the latest task taxonomy update to allocate sufficient time for the response rewrite.

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-03-06T18:07:57.889Z

LLM Judge Step

Issue

Failure to check LLM judges for Response A and B against rubric scoring step

Sometimes the LLM judges for Response A and B near the end of the task do not align with the way you scored those responses early in the task. Watching out for this can help you catch issues.

Approach

- Go back to check score alignment.
- Make corrections if needed.
- Provide a clear justification if you have a reason to disagree with the LLM judge's verdict.

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-03-06T18:07:57.889Z

Working with Generated Rubric Suggestions & Evals

Auto-generated rubrics and evals are plausible but often have subtle issues, and they may not be complete.

Always use your own judgement and ensure that your rubrics are complete and accurate.

The Rubric **needs to be structured in a very specific way**, or the model might learn the wrong things from the rubric.

Consider rubric suggestions and evals as valuable tools, but always decide for yourself.

Rubrics are highly prompt-specific, but...

...the goal is a set of criteria that are **mutually exclusive** (they shouldn't overlap) and **collectively exhaustive** (taken together, they should fully define the ideal response).

MECE can be achieved by following **three principles:**

Atomicity

Each criterion should cover **exactly one** aspect

Specificity

Criteria should **precisely** state what is expected

Self-containment

Each criterion should contain **all the information** needed to evaluate a response.

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-03-06T18:07:57.889Z

Thank You

Outlier

CONFIDENTIAL | ruban.arnold@gmail.com+outlier | 2026-03-06T18:07:57.899Z