Common Errors Webinar

Outlier

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Purpose ▪ Scope ▪ Structure

Purpose

Overview and clarification of common errors in order to improve quality alignment.

Scope

Project-wide common errors for Eunomia only. Any individual account issues should be addressed by reaching out to a QM.

Structure

Common Errors presentation with time for clarification questions.
Please keep your mic muted when not speaking. This helps everyone focus.
Clarification questions may be asked by raising your hand or adding the question to the Zoom chat.

CONFIDENTIAL | ruban.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Rubrics Criteria Objectivity/Specificity

Failing Issue

Objectivity / Specificity

[Fail - Objectivity / Specificity]: Criterion corresponding to explicit requirements is subjective, vague, or unmeasurable.

Examples

❌	"The response should have good formatting"
✅	"The response includes Python code that adheres to PEP 8 indentation standards"
❌	"Code must be optimal"
✅	"The response uses a hash map to achieve O(n) time complexity for the lookup operation, avoiding nested loops."

Write criteria that encompass all requirements needed to fulfill this prompt. 15/15 completed

1 **For allow_multiple: true with a groups mapping, the script selects at most one item per group.**
Code Rubric · Mandatory

2 **The code must check that every group-selected item must be present in that category's items list.**
Code Rubric · Mandatory

3 **After group picks, the script fills remaining slots (up to max_picks) from items not already chosen.**
Code Rubric · Mandatory

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Avoid Generic Rubrics

🛑 Error Type: Generic Rubrics (Boilerplate Criteria)

Definition: Using "copy-paste" criteria that are not tailored to the specific technical task. These are abstract placeholders (like asking for a summary, next steps, or specific formatting) that could apply to almost any prompt regardless of the programming language or problem.

The "Specific to Task" Test: Ask yourself: "If I swapped this Python prompt for a Java or SQL prompt, would this rubric still work perfectly?" If the answer is yes, the rubric is likely too generic and lacks "technical teeth."

Contextual Exceptions (When it's a Non-Fail):

- **Quality of Life (QoL) Threshold:** Generic criteria are acceptable **only** if they are classified as "Quality of Life" and weighted as **"Nice-to-Have."**

Example Case:

- **Prompt:** "How do I split a string in Python using a delimiter?"
- ❌ **Bad Practice (Generic/Boilerplate):**
  - **Criterion 8:** "The response offers follow-up questions or next steps."
  - **Criterion 11:** "The response provides a brief summary or recap."
  - **Criterion 12:** "The response avoids decorative or excessive formatting."
  - **Why it's weak:** These tell us nothing about Python string methods (`.split()`) or delimiters. They are "filler" rubrics.
- ✅ **Good Practice (Task-Specific):**
  - **Criterion 1:** "The response demonstrates the `.split()` method and explains how to pass a specific string delimiter as an argument."
  - **Criterion 2:** "The code correctly handles edge cases, such as splitting a string that does not contain the specified delimiter."
  - **Criterion 3 (Specific QoL):** "The response suggests using `re.split()` if the user needs to split by multiple different characters at once." (This is a "next step" that is actually **specific** to the task).

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Watch out for vague wording

Error Category	Issue Description	Bad Example (Vague/Dependent)	Good Example (Specific/Self-Contained)
Vague Adjectives	Using subjective words like "appropriate," "correctly," or "valid" which an LLM judge cannot define.	"The response suggests a tech stack that is appropriate for the scale."	"The response suggests a tech stack **using SQLite and a single Python script** for a scale of <10k records."
Lack of Ground Truth	Failing to specify the actual technical method required, forcing the judge to guess the "right" way.	"The response provides a technically valid alternative to track.label."	"The response provides an alternative **such as using the videoKind property** to distinguish the stream."
Non-Self-Contained	Requiring the judge to refer back to the prompt or external docs to understand the criteria.	"The model follows the specific instructions mentioned in the prompt."	"The model implements the AES-256 encryption using the **cryptography library as requested.**"
Placeholder Language	Using "filler" phrases that sound professional but lack technical "teeth."	"The response follows industry best practices for error handling."	"The response includes a **try-except block** that catches the specific FileNotFoundError."
Circular Logic	Defining the rubric by the success of the outcome rather than the steps to get there.	"The code works correctly and produces the right output."	"The code uses a list comprehension to filter the data, ensuring the output list contains only integers."
High-Weight Cosmetics	Giving too much importance to formatting (Quality of Life) instead of technical logic.	"The response uses perfectly formatted headers and bold text."	"The response uses Markdown headers to separate the 'Setup' and 'Execution' phases for readability."
QoL rubrics passing All 4 models	The project attempter instructions explicitly say that "Quality-of-Life rubric evaluates the differences between the responses. **A good QoL criterion cannot be passed by all four responses."** QoL rubrics MUST evaluate the differences between responses.	a QoL rubric that all 4 responses pass or all 4 fail.	QoL rubrics MUST evaluate the differences between responses. Check: Ask yourself: "Does this rubric actually help distinguish which model is better?" 🔍

02/13 TLDR: Stop using adjectives or vague words, start using technical nouns and give concrete examples for specific methods. 🎯

Example: Don't say "appropriate tech stack"; say "uses SQLite and Python." ✅
Don't say "uses a valid method"; say "uses the videoKind property." ✅
Don't say "explains clearly"; say "identifies the LZ-String algorithm." ✅

But also watch out for being too prescrptive, you should give multiple examples if the model uses different ways to implement a code, don't limited the model to a single method!!

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Self-contained Criteria

Failing Issue

Self-contained Rubric

[Fail - Self-contained Criteria]: At least one rubric criterion cannot be evaluated without access to prompt, reference text, or external facts.

Not self-contained

"Response identifies the first president of the USA"

"The response addresses the bug mentioned in the prompt

Self-contained

"Response identifies the first president of the USA as George Washington"

"The response addresses the bug where the submit button doesn't work"

CONFIDENTIAL | ruban.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Rubrics should be self-contained!!

The Rubric **needs to be structured in a very specific way**, or the model might learn the wrong things from the rubric.

We can codify the features of a useful rubric as a set of principles.

Rubrics are highly prompt-specific, but...

...the goal is a set of criteria that are **mutually exclusive** (they shouldn't overlap) and **collectively exhaustive** (taken together, they should fully define the ideal response).

MECE can be achieved by following **three principles:**

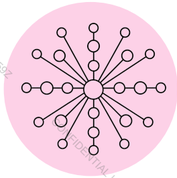

Atomicity

Each criterion should cover **exactly one** aspect

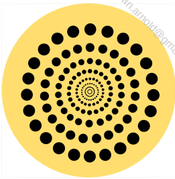

Specificity

Criteria should **precisely** state what is expected

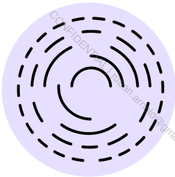

Self-containment

Each criterion should contain **all the information** needed to evaluate a response.

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Rubrics Criteria Objectivity/Specificity

Failing Issue

Overlap / Redundancy

[Fail - Overlap / Redundancy]: At least one criterion is completely redundant.

This double penalizes models and skews ratings.

Redundant Scenario 1: Criteria 1: Response does a, b, c Criteria 2: Response does a, b	**Redundant Scenario 2:** Criteria 1: Response does a, b, c Criteria 2: Response does b, c, d

Examples

Redundant Criteria:

- ❌ "The database schema definition includes an index on the email column of the users table."
- ❌ "The SQL migration script optimizes the users table for frequent lookups by email address."

Non-Redundant Criterion:

- ✅ "The database schema creation script defines a lookup structure on the users.email column (e.g. via a PRIMARY KEY, UNIQUE constraint, or CREATE INDEX statement).

Note: This applies to cases where two criteria independently assess the same elements, not when a single criterion introduces and specifies related requirements."The response follows best code practices by ensuring that each line is under 79 characters" is acceptable).

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Coverage

Failing Issue

Missing Coverage

[Fail - Missing Coverage]: The rubric is missing at least one criterion that checks for an explicit requirement in the prompt or a critical implicit expectation.

Example

The user request requires a Cloudflare Worker with a Cron Trigger that runs every 5 minutes, but the 5-minute timeframe is not specified in the rubric.

CONFIDENTIAL | ruban.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Coverage & Accuracy

Failing Issue

[Fail - Criteria Accuracy]

- Checks for something that contradicts the prompt.
- Contains factual errors or misleading points.
- Not an explicit requirement and doesn't improve the response.

Examples

- ✕ **Misleading and/or factually incorrect** statements like *"The response implements a sorting algorithm that runs in O(nlogn), such as selection sort"*
- ✕ The prompt specifies Python, but the rubric requires JavaScript instead.

CONFIDENTIAL | ruban.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Atomicity

Failing Issue

Atomicity of Rubric Criteria

[Fail - Criteria Atomicity]: Criterion groups unrelated constraints. Results in a rubric item with no clear focus.

It should read as a single coherent instruction, not a dump of requirements.

Example

Non-Atomic Criterion:

❌ "The function wraps the API call in a try/catch block to handle network errors and logs the error message to an external monitoring service like Sentry.

Atomic Criterion:

✅ The response includes a try/catch block around the API request logic to handle network errors.

✅ The catch block invokes a logging function or service to record error details.

CONFIDENTIAL | ruban.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Missing Criteria

Failing Issue

Missing Rubric Criteria

[Fail - Missing Criteria]
Criteria that are objectively 100% essential to meet the fundamental requirements of the prompt are missing, preventing the response from successfully fulfilling the core task if not included.

Example

Prompt:
*Write a Python function using the requests library to fetch **all** user records from the /api/users endpoint, handling **pagination** via the next_page_url field in the JSON response.*

✕ No criterion requires the code to **loop or recurse**, so code that returns only the first page of users could potentially pass the rubrics while failing to fulfill the core request.

CONFIDENTIAL | ruban.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Counterproductive Criteria

Failing Issue

Counterproductive Criteria

[Fail - Counterproductive Criteria]
At least one criterion penalizes good responses / rewards bad responses such that the inclusion of this criteria would make the response objectively worse.

Example

- ✕ A rubric imposes a technical constraint (like specific text-based class names) that is **incompatible** with the subject's fundamental architecture (e.g., randomized/hashed builds), effectively forcing a brittle solution.

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Formatting

Failing Issue

Formatting of Rubric Criteria

[Fail - Multiple Formatting Issues]: 2+ criteria have one of the following formatting issues:

- Written in past tense
- Written as a question
- Starts with "The model"

Example

Bad Rubric Criteria Formatting:

❌ "The model used a context manager (e.g., with open(...) or Path(...).open()) for all file I/O operations."

❌ "Did the response invoke a context manager (e.g., with open(...) or Path(...).open()) for file I/O operations?"

Correctly Formatted Criterion:

✅ "The script utilizes a context manager (e.g., with open(...) or Path(...).open()) for all file I/O operations."

CONFIDENTIAL | ruban.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Rubric Category definitions

Code

→ Evaluates the **correctness, safety,** and **intent** of the code

- Focus on intent, not exact syntax
- Avoid requiring specific methods, libraries, or line-by-line implementations unless explicitly needed
- Accept alternative valid approaches if they meet the goal

Instruction Following (IF)

→ Ensures the model correctly follows the user's **explicit and implicit instructions**

- Check that all parts of the prompt are fully addressed
- No additions, omissions, or scope drift
- No reinterpreting or ignoring user requirements
- **Instruction Following (IF) rubrics must be mandatory.**

Quality of Life (QoL)

→ Evaluates usefulness and clarity across three sub-dimensions:

- **Helpfulness** – relevance, usefulness, guidance, and solving the user's need
- **Style & Clarity** – readable code, clear structure, logical flow, visual organization
- **Verbosity** – concise yet complete; no rambling, no missing essential information

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Criteria Weights: Mandatory vs. Nice to Have

Failing Issue

[Fail - Criteria Weights Major Issue]:
>20% of the criteria weights are objectively wrong (Mandatory vs Nice to Have).

Mandatory

- Covers **explicit** or **implicit** requirements from the prompt.
- Represents something the response **must include** to be an acceptable response.

Nice to Have

- Improves response quality but is **not required** to fulfill the task.
- Reflects enhancements that **objectively** improve usefulness or clarity.

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Vague or Irrelevant Criteria & Rewrite Issues

Vague / Generic

[Fail - Vague Criteria]

The criterion is too vague or abstract to be meaningfully evaluated, making it impossible to determine whether it has been met.

[Fail - Generic Criteria]

Non-QoL criterion is too generic (could be applied to virtually any prompt).

Irrelevant Rubric

[Fail- Irrelevant Criteria]

Criterion defined does not make the response objectively better OR it makes the response objectively worse.

Code Rewrite

[Fail - Code Rewrite Requirements]

The rewrite does NOT pass all the criteria.

If a rewrite fails any rubric, the task fails and the highest quality score it can receive is 2.

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z

Thank You

Outlier

CONFIDENTIAL | ruben.arnold@gmail.com+outlier | 2026-02-25T01:18:45.859Z